

Introduction to Transfer Learning

Thien Tran

tran0096@e.ntu.edu.sg

MLDA@EEE

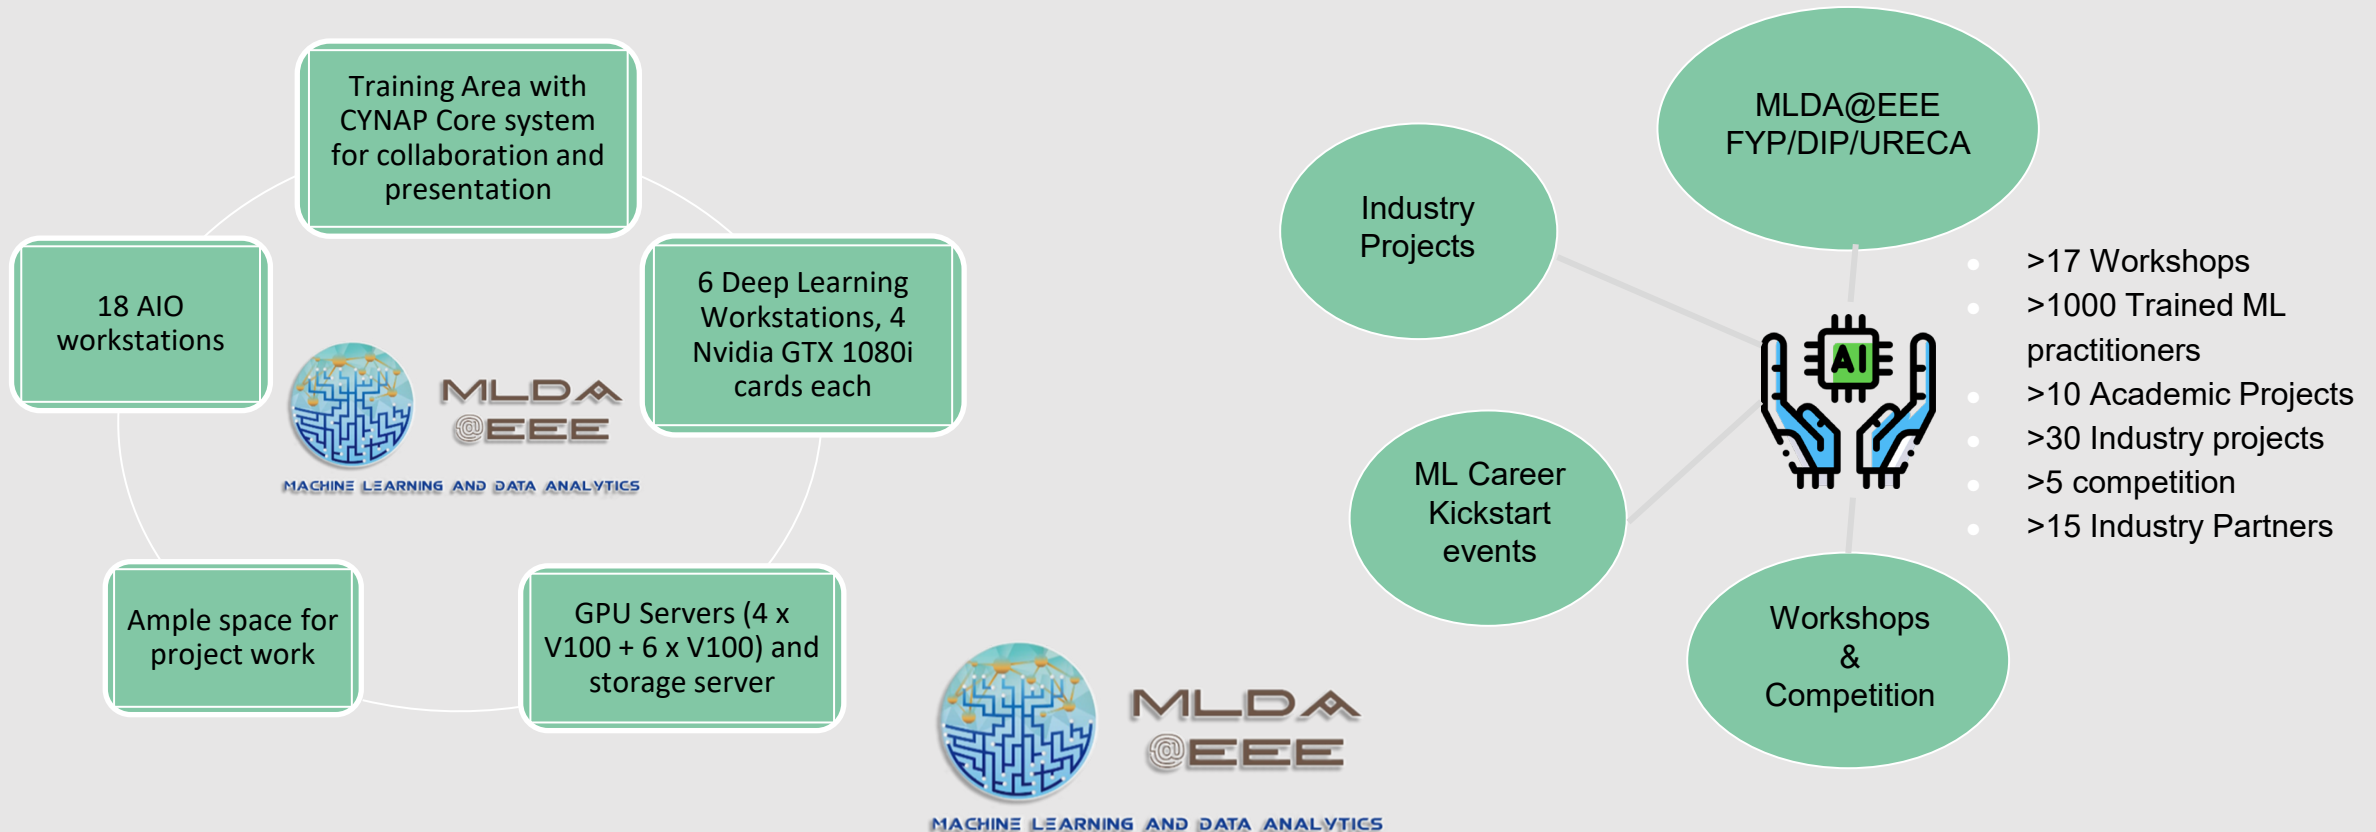


MLDA
@EEE

MACHINE LEARNING AND DATA ANALYTICS

Our Mission

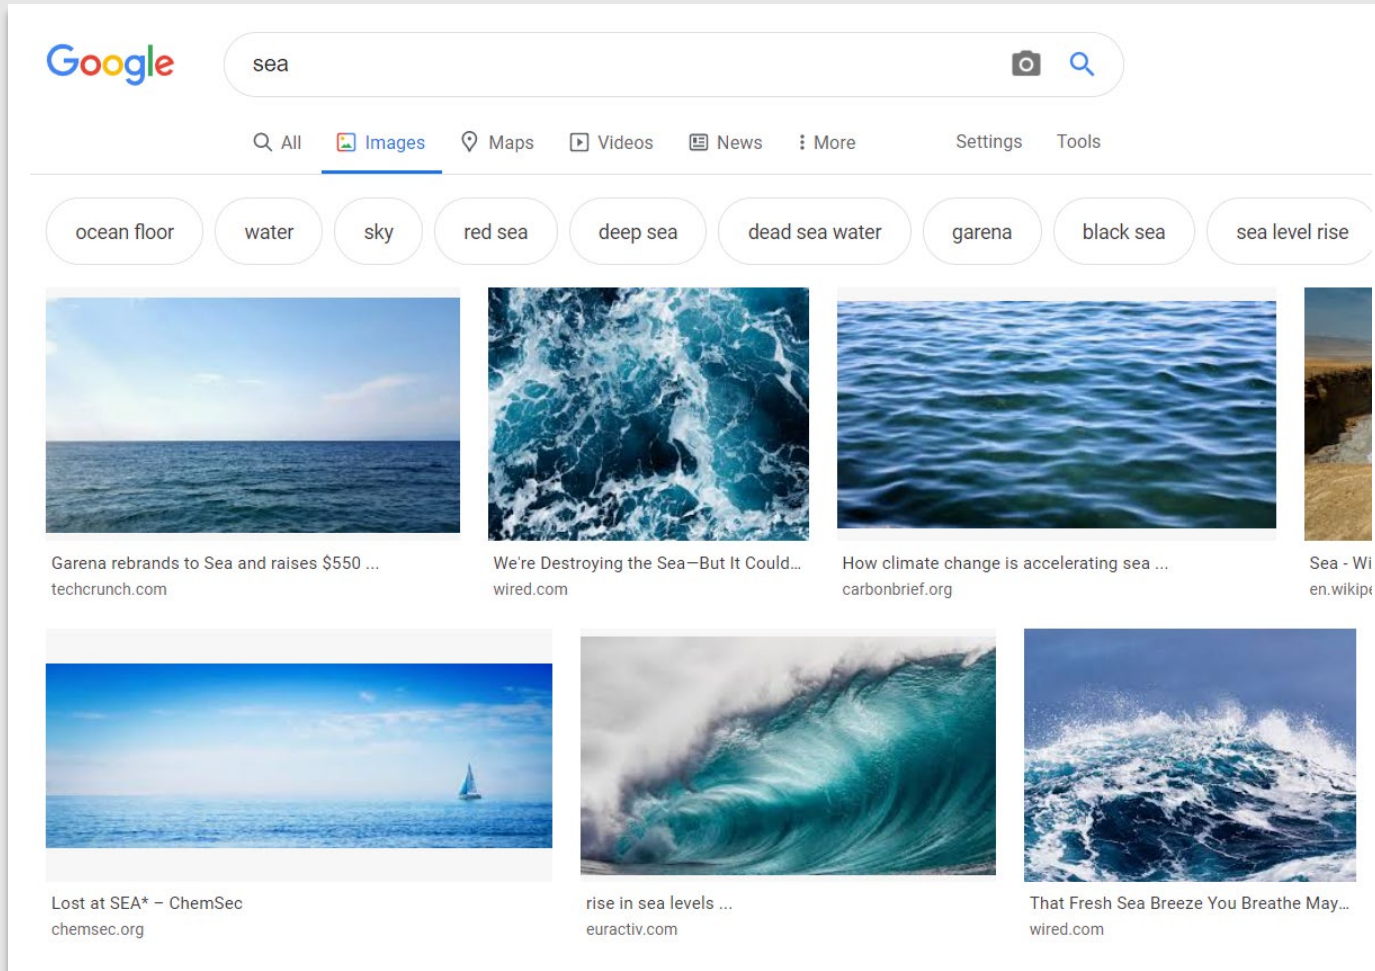
Provide an integrated platform for EEE/IEM students to learn and implement Machine Learning, Data Science & AI, as well as facilitate connections with the industry.



Agenda

- Theory of Transfer learning
- Transfer Learning workflow
- Hands-on: implement Transfer Learning with TensorFlow Keras

Image auto-tagging



You have a collection of a lot of images

You want to view your **'sea'** images
And potentially all other types of
landscape images also

You can label some of them manually
but cannot label the whole collection

















MACHINE LEARNING TASK

Image
classification

CONSTRAINTS

Small, limited
dataset

Classify landscape images



Intel Image Classification

<https://www.kaggle.com/puneet6060/intel-image-classification>

What is Transfer Learning

- Apply knowledge learned from **one task** to another **related task**

LAKE CLASSIFIER



NTU Yunnan Garden

<https://media.ntu.edu.sg/NewsReleases/Pages/newsdetail.aspx?news=a9f88bb3-c20b-4dea-a1e2-7eb65a4c8d9b>



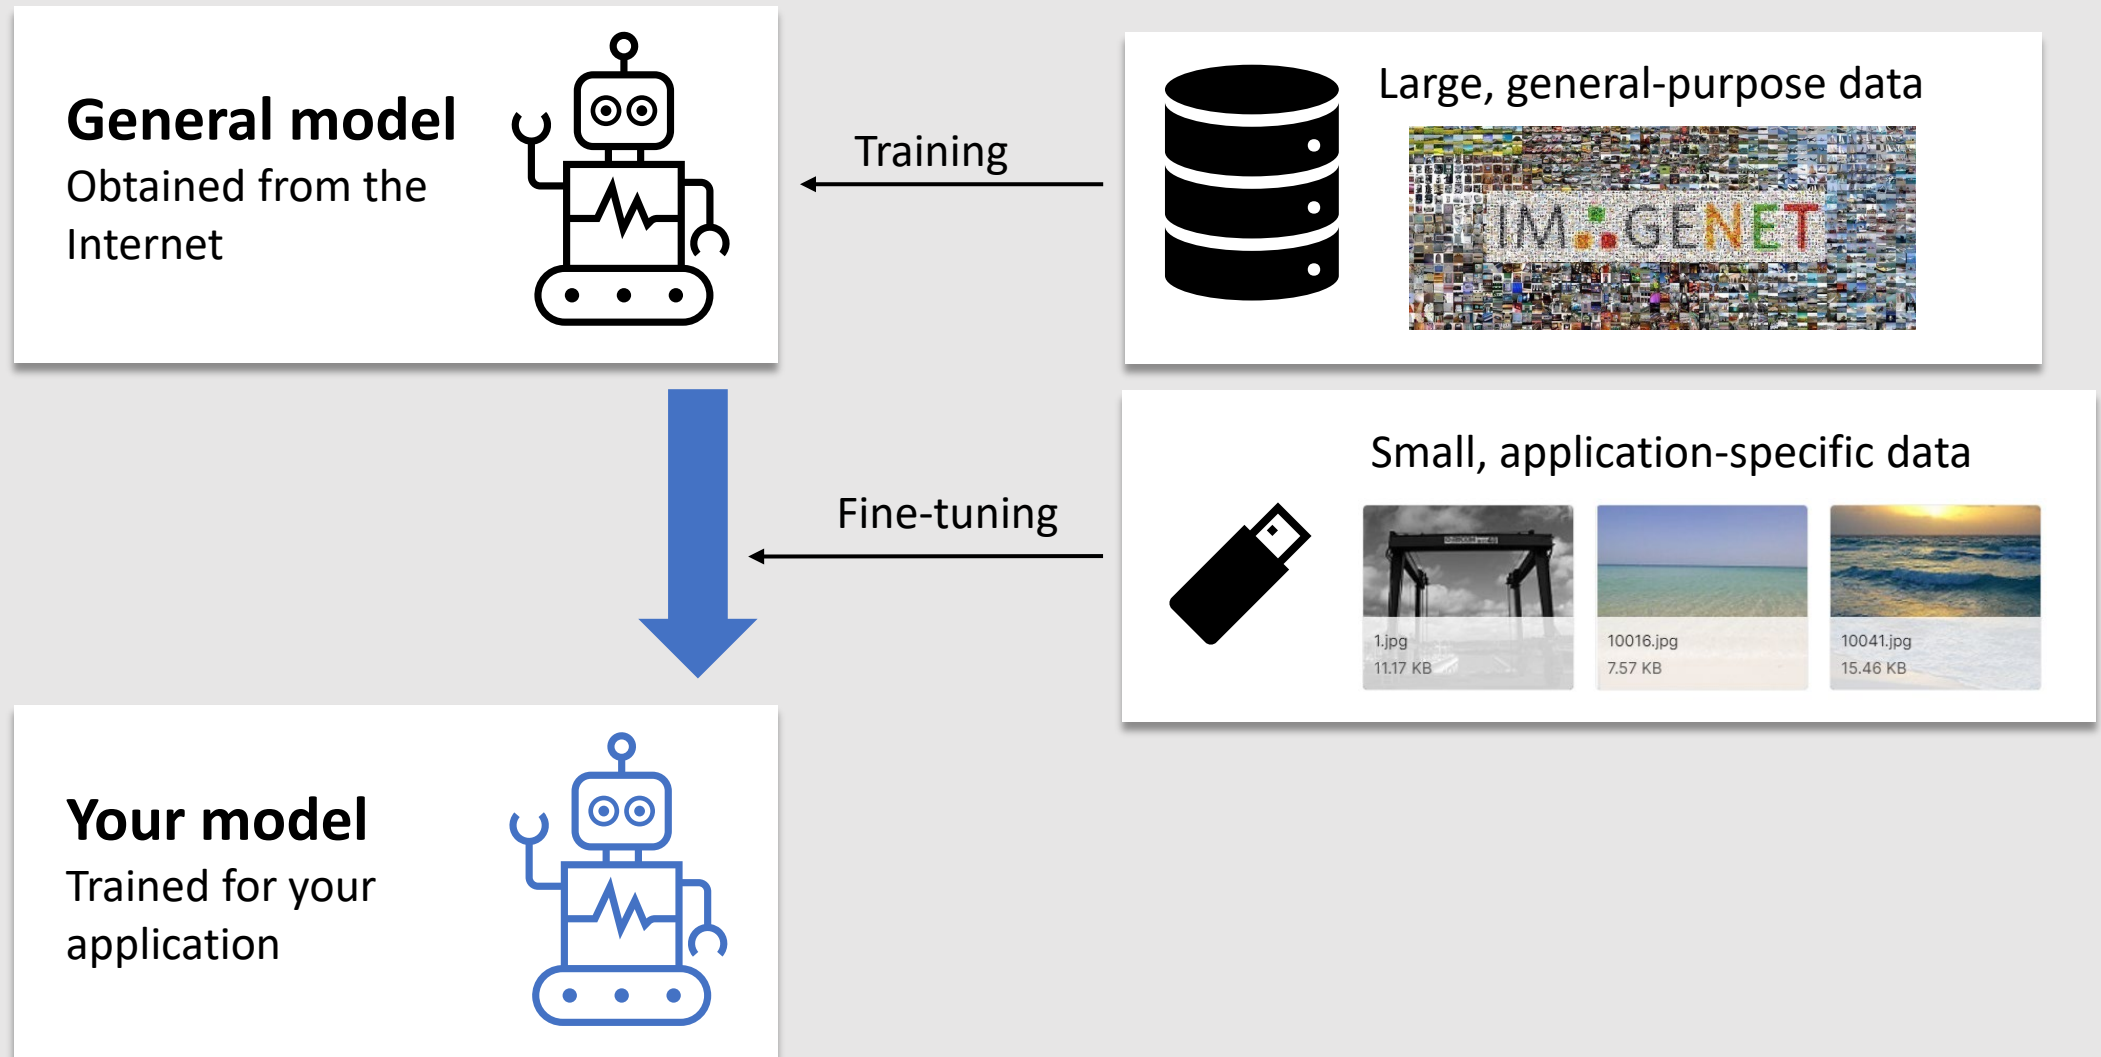
SEA CLASSIFIER



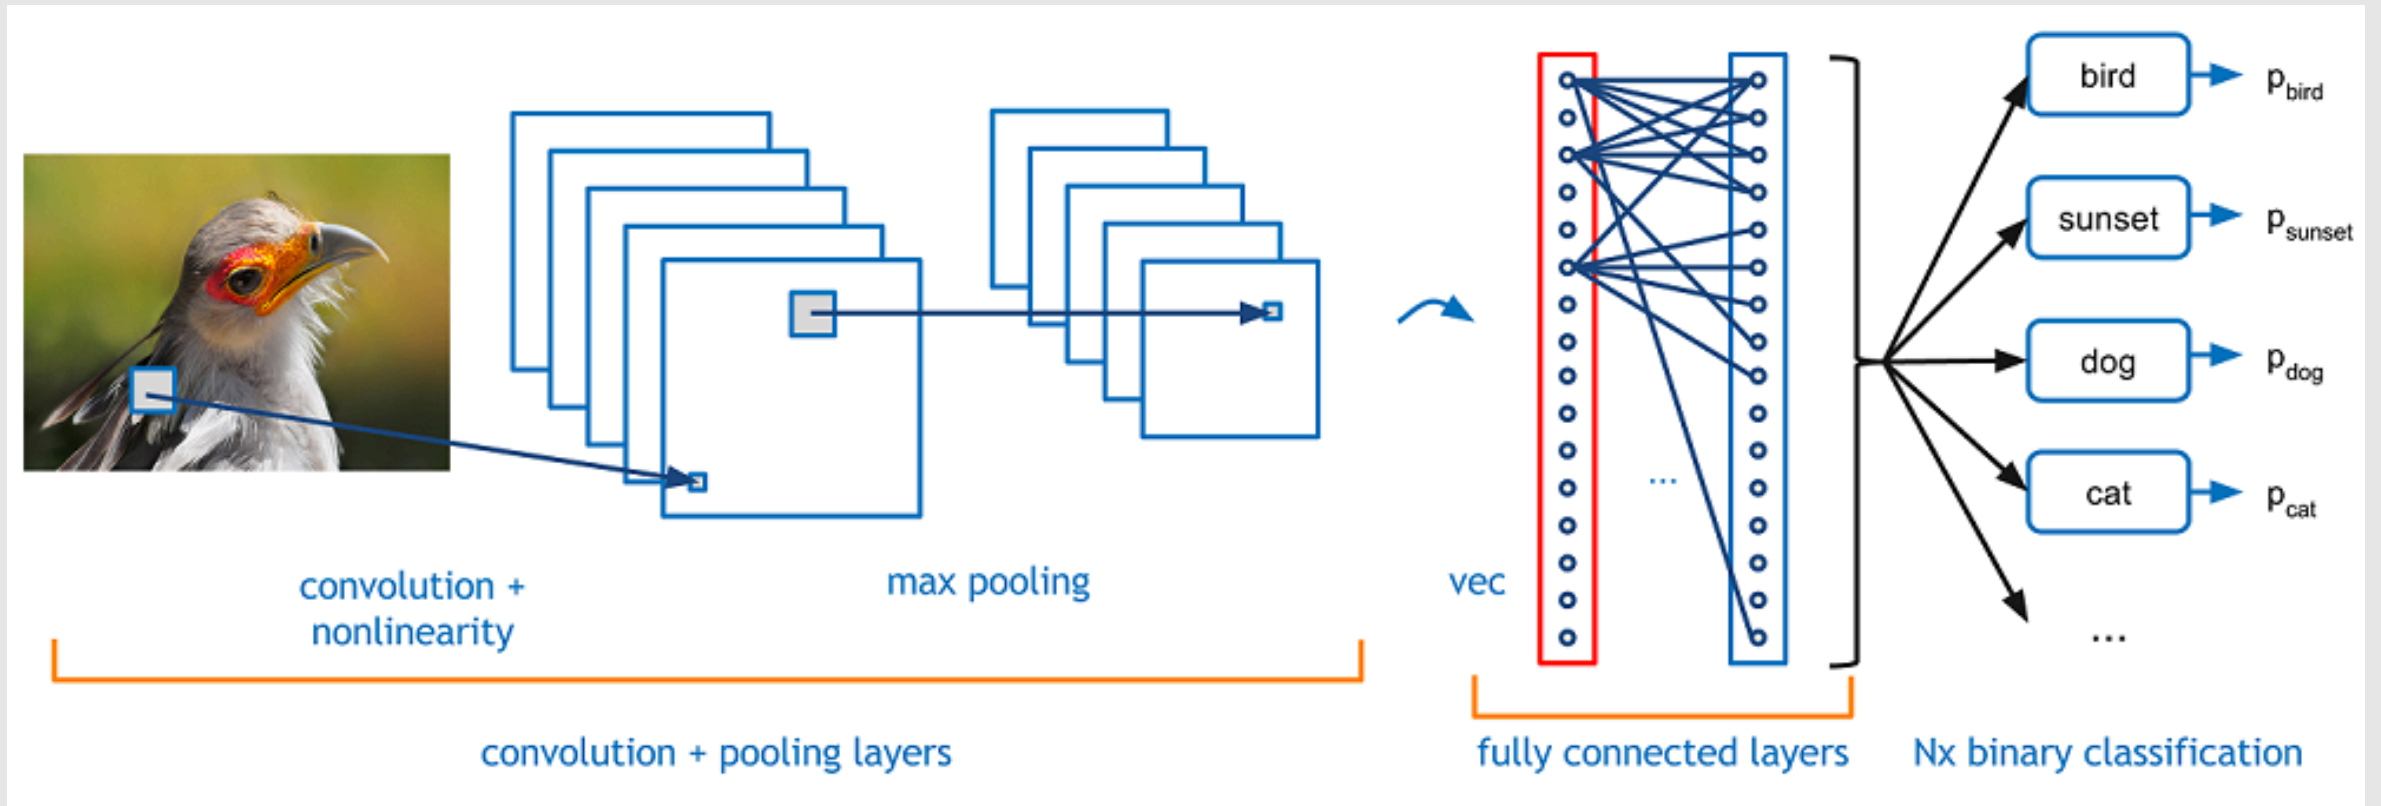
'sea' from Intel Image Classification

<https://www.kaggle.com/puneet6060/intel-image-classification>

What is Transfer Learning



Transfer Learning – Learning the features





Extract Image features

Features summary


Transfer Learning – Black box




Input – Rank 3 Tensor



Black box – Pre-trained model




Output - Features vector



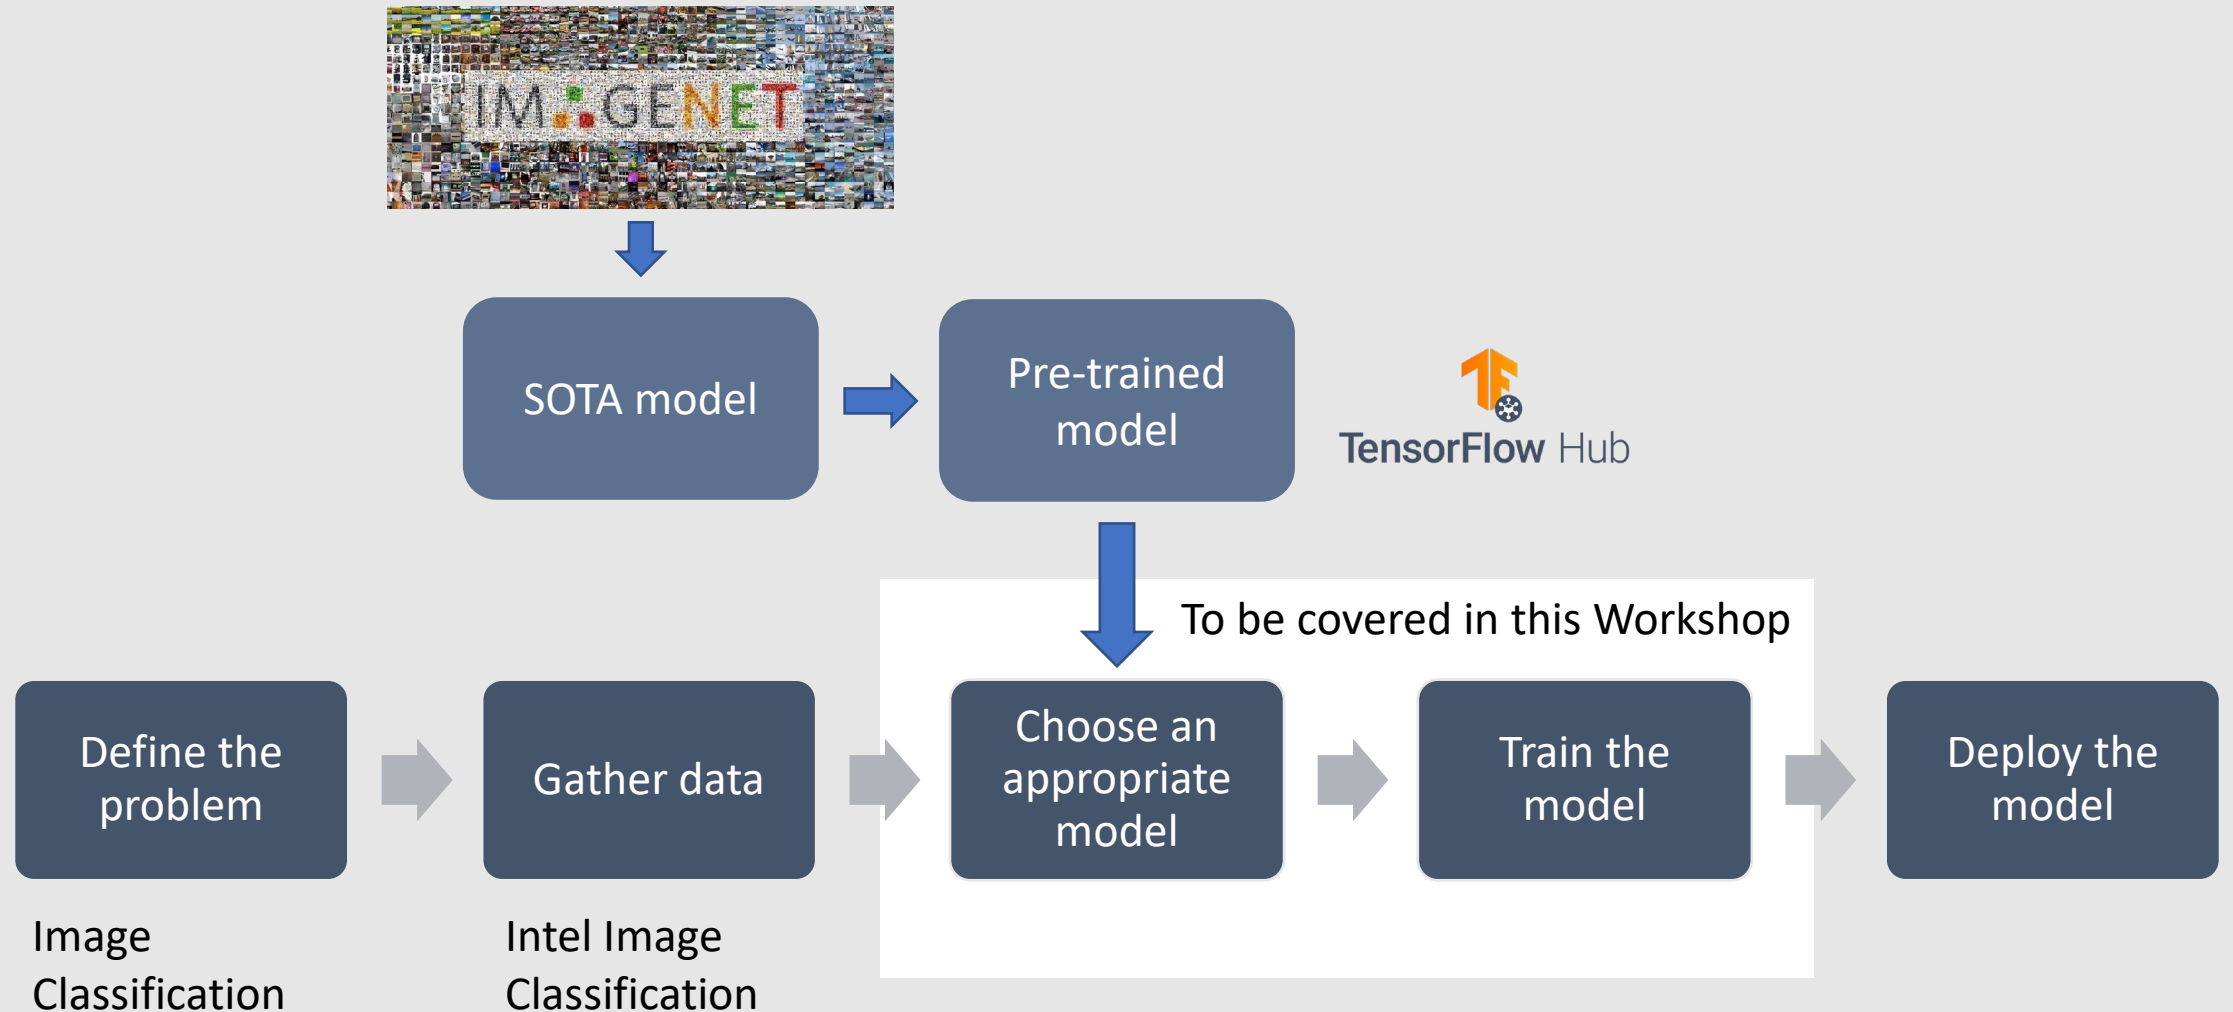
Classification head

Transfer Learning – Black box



Features vectors of Intel Image
Classification dataset in 3D space
(Generated from Google's BiT-M model)

Sample Transfer Learning Workflow



Upstream Training - ImageNet

- ImageNet: a database of images for visual object recognition research
 - 14 million images, hand-annotated
 - 20,000 categories (classes)
 - The standard dataset for evaluating neural network architecture in research
 - <http://image-net.org/explore>



Upstream Training - ImageNet

Geological formation, formation

(geology) the geological features of the earth

1808 pictures 86.24% Popularity Percentile Wordnet IDs

Numbers in brackets: (the number of synsets in the subtree).

- ImageNet 2011 Fall Release (32326)
 - plant, flora, plant life (4486)
 - geological formation, formation (17)
 - aquifer (0)
 - beach (1)
 - cave (3)
 - cliff, drop, drop-off (2)
 - delta (0)
 - diapir (0)
 - folium (0)
 - foreshore (0)
 - ice mass (10)
 - lakefront (0)
 - massif (0)
 - monocline (0)
 - mouth (0)
 - natural depression, depression (0)
 - natural elevation, elevation (41)
 - oceanfront (0)
 - range, mountain range, range of mountains (0)
 - relict (0)
 - ridge, ridgeline (2)
 - ridge (0)
 - shore (7)
 - slope, incline, side (17)
 - spring, fountain, outflow, outpouring (0)
 - talus, scree (0)
 - vein, mineral vein (1)
 - volcanic crater, crater (2)
 - wall (0)

Treemap Visualization Images of the Synset Downloads


ImageNet 2011 Fall Release > Geological formation, formation

Natural 	Slope 	Shore
	Ice 	Water
	Vein 	Delta
	Massif 	Foreshore
	Mouth 	Talus
	Lakefront 	Volcanic
	Wall 	Beach
	Monocline 	Diapir
	Oceanfront 	Cliff
	Aquifer 	Cave
		Spring
		Ridge

SOTA models


Model	Top-1 accuracy
AlexNet (2012)	63.3%
VGG-19 (2015)	74.5%
ResNet-152 (2016)	77.8%
ResNeXt-101 (2017)	80.9%
EfficientNet-B7 (2019)	84.4%
FixEfficientNet-L2 (2020)	88.5%

Image Classification on ImageNet




Source: <https://paperswithcode.com/sota/image-classification-on-imagenet>

Choose a model



Accuracy vs Throughput Trade-off



Model size

Define the
problem

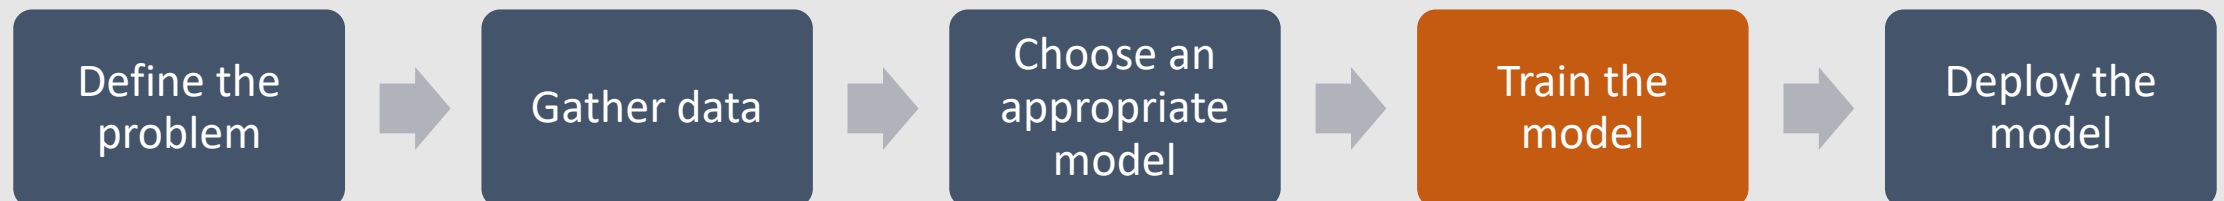

Gather data

Choose an
appropriate
model

Train the
model


Deploy the
model

Prepare the model




Classification layer and Fine-tuning

Train the classification layer only



Train the whole network

Train selectively the last few layers



(Will not be covered during Hands-on)

Define the
problem

Gather data

Choose an
appropriate
model

Train the
model


Deploy the
model

TensorFlow and Keras

TensorFlow's **high-level APIs** are based on the Keras API standard for defining and training neural networks.

Keras enables **fast** prototyping, state-of-the-art research, and production—all with **user-friendly APIs**.

TensorFlow Architecture



TensorFlow and Keras

```
import tensorflow as tf
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

[Run code now](#)[Try in Google's interactive notebook](#)

```
class MyModel(tf.keras.Model):
    def __init__(self):
        super(MyModel, self).__init__()
        self.conv1 = Conv2D(32, 3, activation='relu')
        self.flatten = Flatten()
        self.d1 = Dense(128, activation='relu')
        self.d2 = Dense(10, activation='softmax')


    def call(self, x):
        x = self.conv1(x)
        x = self.flatten(x)
        x = self.d1(x)
        return self.d2(x)

model = MyModel()

with tf.GradientTape() as tape:
    logits = model(images)
    loss_value = loss(logits, labels)
    grads = tape.gradient(loss_value, model.trainable_variables)
    optimizer.apply_gradients(zip(grads, model.trainable_variables))
```

[Run code now](#)[Try in Google's interactive notebook](#)

Hands-on session




<https://github.com/MLDA-NTU/Transfer-Learning-DL2020>

Transfer Learning for NLP tasks

Upstream training with Wikipedia data



Word embeddings



Word2Vec 10K

<https://projector.tensorflow.org>

Language models

GPT-3
BERT
Turing-NLG
RoBERTa

Can be obtained from



Hugging Face

Feedback



https://docs.google.com/forms/d/e/1FAIpQLSdxmE7yQsHjNOZVVY654AQufyxJHMF-G_ytjHg4ep7ZBiL5Mw/viewform

Thien Tran
tran0096@e.ntu.edu.sg

MLDA@EEE